

L3 Exercise 2 - IaC

July 4, 2021

1 Exercise 2: Creating Redshift Cluster using the AWS python SDK

1.1 An example of Infrastructure-as-code

```
In [67]: import pandas as pd
import boto3
import json
```

2 STEP 0: Make sure you have an AWS secret and access key

- Create a new IAM user in your AWS account
- Give it AdministratorAccess, From Attach existing policies directly Tab
- Take note of the access key and secret
- Edit the file dwh.cfg in the same folder as this notebook and fill [AWS] KEY= YOUR_AWS_KEY SECRET= YOUR_AWS_SECRET

3 Load DWH Params from a file

```
In [68]: import configparser
config = configparser.ConfigParser()
config.read_file(open('dwh.cfg'))

KEY = config.get('AWS', 'KEY')
SECRET = config.get('AWS', 'SECRET')

DWH_CLUSTER_TYPE = config.get("DWH", "DWH_CLUSTER_TYPE")
DWH_NUM_NODES = config.get("DWH", "DWH_NUM_NODES")
DWH_NODE_TYPE = config.get("DWH", "DWH_NODE_TYPE")

DWH_CLUSTER_IDENTIFIER = config.get("DWH", "DWH_CLUSTER_IDENTIFIER")
DWH_DB = config.get("DWH", "DWH_DB")
DWH_DB_USER = config.get("DWH", "DWH_DB_USER")
DWH_DB_PASSWORD = config.get("DWH", "DWH_DB_PASSWORD")
DWH_PORT = config.get("DWH", "DWH_PORT")

DWH_IAM_ROLE_NAME = config.get("DWH", "DWH_IAM_ROLE_NAME")
```

```
(DWH_DB_USER, DWH_DB_PASSWORD, DWH_DB)
```

```
pd.DataFrame({"Param":
              ["DWH_CLUSTER_TYPE", "DWH_NUM_NODES", "DWH_NODE_TYPE", "DWH_CLUSTER_IDENTIFI",
              "Value":
              [DWH_CLUSTER_TYPE, DWH_NUM_NODES, DWH_NODE_TYPE, DWH_CLUSTER_IDENTIFI
              ]})
```

```
Out[68]:
```

	Param	Value
0	DWH_CLUSTER_TYPE	multi-node
1	DWH_NUM_NODES	4
2	DWH_NODE_TYPE	dc2.large
3	DWH_CLUSTER_IDENTIFIER	dwhCluster
4	DWH_DB	dwh
5	DWH_DB_USER	dwhuser
6	DWH_DB_PASSWORD	PasswOrd
7	DWH_PORT	5439
8	DWH_IAM_ROLE_NAME	dwhRole

3.1 Create clients for EC2, S3, IAM, and Redshift

```
In [69]: import boto3
```

```
ec2 =
```

```
s3 =
```

```
iam =
```

```
redshift =
```

3.2 Check out the sample data sources on S3

```
In [73]: sampleDbBucket = s3.Bucket("awssampleduswest2")
```

```
# TODO: Iterate over bucket objects starting with "ssbgz" and print
```

```
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/customer0002_part_00.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/dwdate.tbl.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/lineorder0000_part_00.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/lineorder0001_part_00.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/lineorder0002_part_00.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/lineorder0003_part_00.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/lineorder0004_part_00.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/lineorder0005_part_00.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/lineorder0006_part_00.gz')
```

```
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/lineorder0007_part_00.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/part0000_part_00.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/part0001_part_00.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/part0002_part_00.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/part0003_part_00.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/supplier.tbl_0000_part_00.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/supplier0001_part_00.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/supplier0002_part_00.gz')
s3.ObjectSummary(bucket_name='awssampleduswest2', key='ssbgz/supplier0003_part_00.gz')
```

3.3 STEP 1: IAM ROLE

- Create an IAM Role that makes Redshift able to access S3 bucket (ReadOnly)

```
In [ ]: # TODO: Create the IAM role
try:
    print('1.1 Creating a new IAM Role')
    dwhRole =

except Exception as e:
    print(e)

In [ ]: # TODO: Attach Policy
print('1.2 Attaching Policy')

In [82]: # TODO: Get and print the IAM role ARN
print('1.3 Get the IAM role ARN')
roleArn =

print(roleArn)
```

1.1 Creating a new IAM Role

An error occurred (EntityAlreadyExists) when calling the CreateRole operation: Role with name dw

1.2 Attaching Policy

1.3 Get the IAM role ARN

arn:aws:iam::988332130976:role/dwhRole

3.4 STEP 2: Redshift Cluster

- Create a RedShift Cluster
- For complete arguments to create_cluster, see [docs](#)

```
In [83]: try:
    response = redshift.create_cluster(
        # TODO: add parameters for hardware
```

```

        # TODO: add parameters for identifiers & credentials

        # TODO: add parameter for role (to allow s3 access)

    )
except Exception as e:
    print(e)

```

An error occurred (ClusterAlreadyExists) when calling the CreateCluster operation: Cluster already exists

3.5 2.1 Describe the cluster to see its status

- run this block several times until the cluster status becomes Available

```

In [77]: def prettyRedshiftProps(props):
    pd.set_option('display.max_colwidth', -1)
    keysToShow = ["ClusterIdentifier", "NodeType", "ClusterStatus", "MasterUsername", "DBName", "Endpoint", "VpcId", "NumberOfNodes"]
    x = [(k, v) for k,v in props.items() if k in keysToShow]
    return pd.DataFrame(data=x, columns=["Key", "Value"])

myClusterProps = redshift.describe_clusters(ClusterIdentifier=DWH_CLUSTER_IDENTIFIER)['ClusterProperties']
prettyRedshiftProps(myClusterProps)

```

```

Out[77]:
      Key \
0  ClusterIdentifier
1  NodeType
2  ClusterStatus
3  MasterUsername
4  DBName
5  Endpoint
6  VpcId
7  NumberOfNodes

      Value
0  dwhcluster
1  dc2.large
2  available
3  dwhuser
4  dwh
5  {'Address': 'dwhcluster.csmamz5zxmle.us-west-2.redshift.amazonaws.com', 'Port': 5439}
6  vpc-54d40a2c
7  4

```

2.2 Take note of the cluster endpoint and role ARN

DO NOT RUN THIS unless the cluster status becomes "Available"

```
In [78]: DWH_ENDPOINT = myClusterProps['Endpoint']['Address']
DWH_ROLE_ARN = myClusterProps['IamRoles'][0]['IamRoleArn']
print("DWH_ENDPOINT :: ", endpoint)
print("DWH_ROLE_ARN :: ", roleArn)

DWH_ENDPOINT :: dwhcluster.csmamz5zxmle.us-west-2.redshift.amazonaws.com
DWH_ROLE_ARN :: arn:aws:iam::988332130976:role/dwhRole
```

3.6 STEP 3: Open an incoming TCP port to access the cluster endpoint

```
In [84]: try:
    vpc = ec2.Vpc(id=myClusterProps['VpcId'])
    defaultSg = list(vpc.security_groups.all())[0]
    print(defaultSg)

    defaultSg.authorize_ingress(
        GroupName= , # TODO: fill out
        CidrIp='', # TODO: fill out
        IpProtocol='', # TODO: fill out
        FromPort=int(DWH_PORT),
        ToPort=int(DWH_PORT)
    )
except Exception as e:
    print(e)

ec2.SecurityGroup(id='sg-d6161da0')
An error occurred (InvalidPermission.Duplicate) when calling the AuthorizeSecurityGroupIngress o
```

3.7 STEP 4: Make sure you can connect to the cluster

```
In [80]: %load_ext sql

The sql extension is already loaded. To reload it, use:
%reload_ext sql

In [81]: conn_string="postgresql://{user}:{password}@{host}:{port}/{db}".format(DWH_DB_USER, DWH_DB_PASSWORD, DWH_ENDPOINT, DWH_DB_NAME)
print(conn_string)
%sql $conn_string

postgresql://dwhuser:PasswOrd@dwhcluster.csmamz5zxmle.us-west-2.redshift.amazonaws.com:5439/dwh

Out[81]: 'Connected: dwhuser@dwh'
```

3.8 STEP 5: Clean up your resources

DO NOT RUN THIS UNLESS YOU ARE SURE We will be using these resources in the next exercises

```
In [85]: ##### CAREFUL!!
```

```
##-- Uncomment & run to delete the created resources
```

```
redshift.delete_cluster( ClusterIdentifier=DWH_CLUSTER_IDENTIFIER, SkipFinalClusterSnapshot=True)
##### CAREFUL!!
```

```
Out[85]: {'Cluster': {'AllowVersionUpgrade': True,
  'AutomatedSnapshotRetentionPeriod': 1,
  'AvailabilityZone': 'us-west-2b',
  'ClusterCreateTime': datetime.datetime(2019, 2, 16, 6, 21, 30, 630000, tzinfo=tzutc()),
  'ClusterIdentifier': 'dwhcluster',
  'ClusterParameterGroups': [{'ParameterApplyStatus': 'in-sync',
    'ParameterGroupName': 'default.redshift-1.0'}],
  'ClusterSecurityGroups': [],
  'ClusterStatus': 'deleting',
  'ClusterSubnetGroupName': 'default',
  'ClusterVersion': '1.0',
  'DBName': 'dwh',
  'Encrypted': False,
  'Endpoint': {'Address': 'dwhcluster.csmamz5zxmle.us-west-2.redshift.amazonaws.com',
    'Port': 5439},
  'EnhancedVpcRouting': False,
  'IamRoles': [{'ApplyStatus': 'in-sync',
    'IamRoleArn': 'arn:aws:iam::988332130976:role/dwhRole'}],
  'MasterUsername': 'dwhuser',
  'NodeType': 'dc2.large',
  'NumberOfNodes': 4,
  'PendingModifiedValues': {},
  'PreferredMaintenanceWindow': 'fri:10:30-fri:11:00',
  'PubliclyAccessible': True,
  'Tags': [],
  'VpcId': 'vpc-54d40a2c',
  'VpcSecurityGroups': []},
  'ResponseMetadata': {'HTTPHeaders': {'content-length': '2041',
    'content-type': 'text/xml',
    'date': 'Sat, 16 Feb 2019 07:13:32 GMT',
    'x-amzn-requestid': '5e58b2d8-31ba-11e9-b19b-0945d449b0a9'},
    'HTTPStatusCode': 200,
    'RequestId': '5e58b2d8-31ba-11e9-b19b-0945d449b0a9',
    'RetryAttempts': 0}}
```

- run this block several times until the cluster really deleted

```
In [86]: myClusterProps = redshift.describe_clusters(ClusterIdentifier=DWH_CLUSTER_IDENTIFIER)
prettyRedshiftProps(myClusterProps)
```

```

Out[86]:                                     Key \
0 ClusterIdentifier
1 NodeType
2 ClusterStatus
3 MasterUsername
4 DBName
5 Endpoint
6 VpcId
7 NumberOfNodes

0 dwhcluster
1 dc2.large
2 deleting
3 dwhuser
4 dwh
5 {'Address': 'dwhcluster.csmamz5zxmle.us-west-2.redshift.amazonaws.com', 'Port': 5439}
6 vpc-54d40a2c
7 4

```

```

In [87]: ##### CAREFUL!!
        #-- Uncomment & run to delete the created resources
        iam.detach_role_policy(RoleName=DWH_IAM_ROLE_NAME, PolicyArn="arn:aws:iam::aws:policy/AWS_IAM_Authentication")
        iam.delete_role(RoleName=DWH_IAM_ROLE_NAME)
        ##### CAREFUL!!

```

```

Out[87]: {'ResponseMetadata': {'HTTPHeaders': {'content-length': '200',
        'content-type': 'text/xml',
        'date': 'Sat, 16 Feb 2019 07:13:50 GMT',
        'x-amzn-requestid': '694f8d91-31ba-11e9-9438-d3ce9c613ef8'},
        'HTTPStatusCode': 200,
        'RequestId': '694f8d91-31ba-11e9-9438-d3ce9c613ef8',
        'RetryAttempts': 0}}

```

```

In [ ]:

```